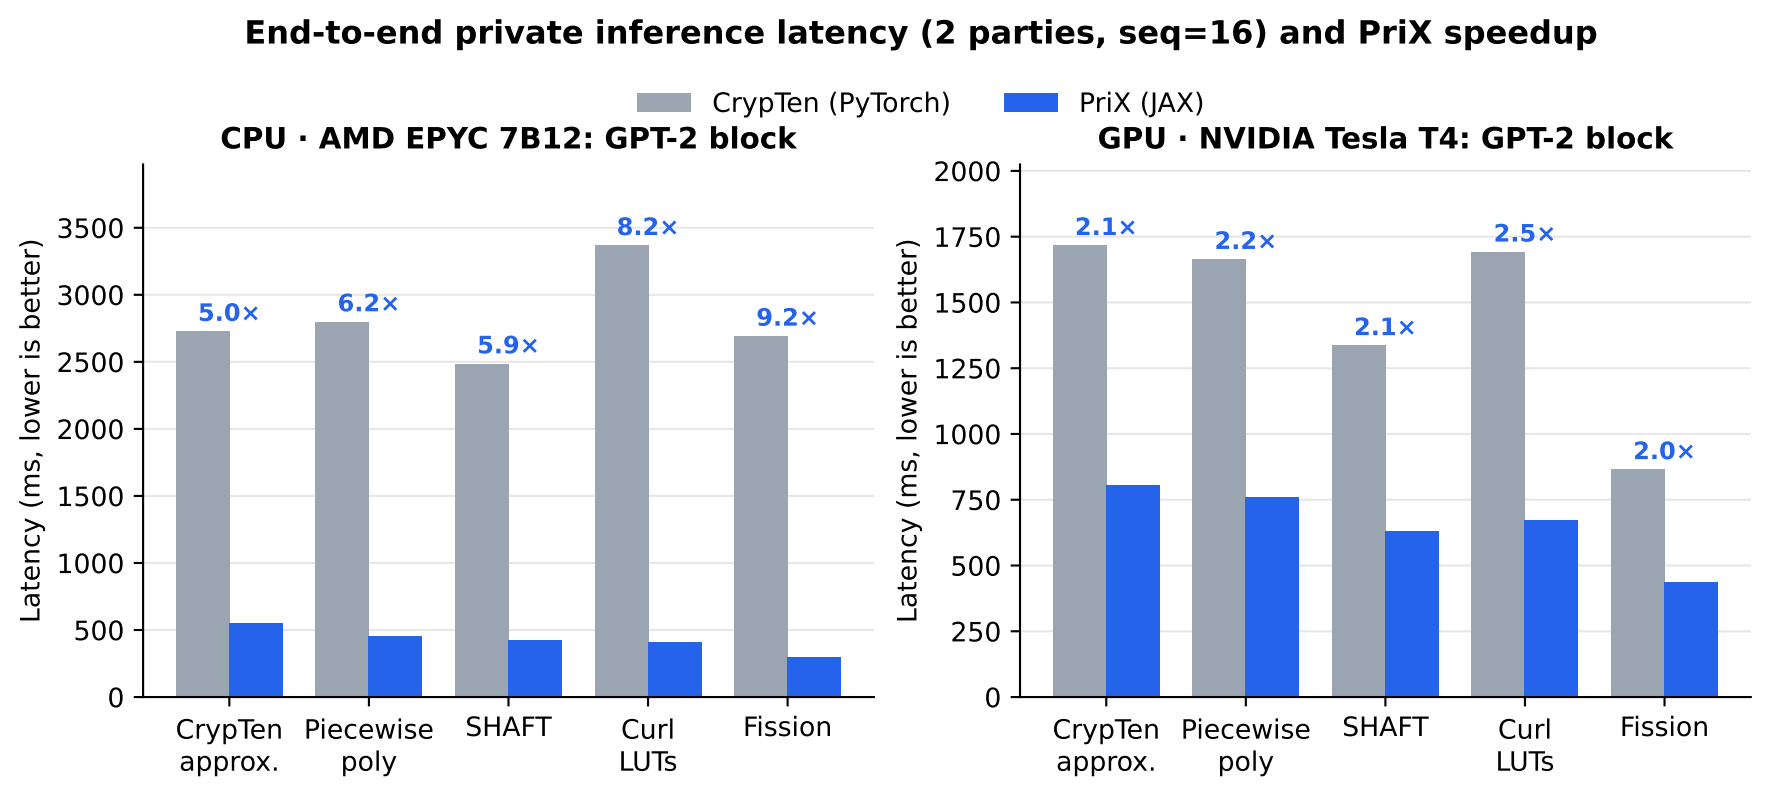
End-to-end private inference latency (2 parties, seq=16) and PriX speedup
CrypTen (PyTorch)
PriX (JAX)
CPU · AMD EPYC 7B12: GPT-2 block
GPU · NVIDIA Tesla T4: GPT-2 block
Latency (ms, lower is better)
Latency (ms, lower is better)
0
500
1000
1500
2000
2500
3000
3500
0
250
500
750
1000
1250
1500
1750
2000
5.0×
6.2×
5.9×
8.2×
9.2×
2.1×
2.2×
2.1×
2.5×
2.0×
CrypTen approx.
Piecewise poly
SHAFT
Curl LUTs
Fission
CrypTen approx.
Piecewise poly
SHAFT
Curl LUTs
Fission