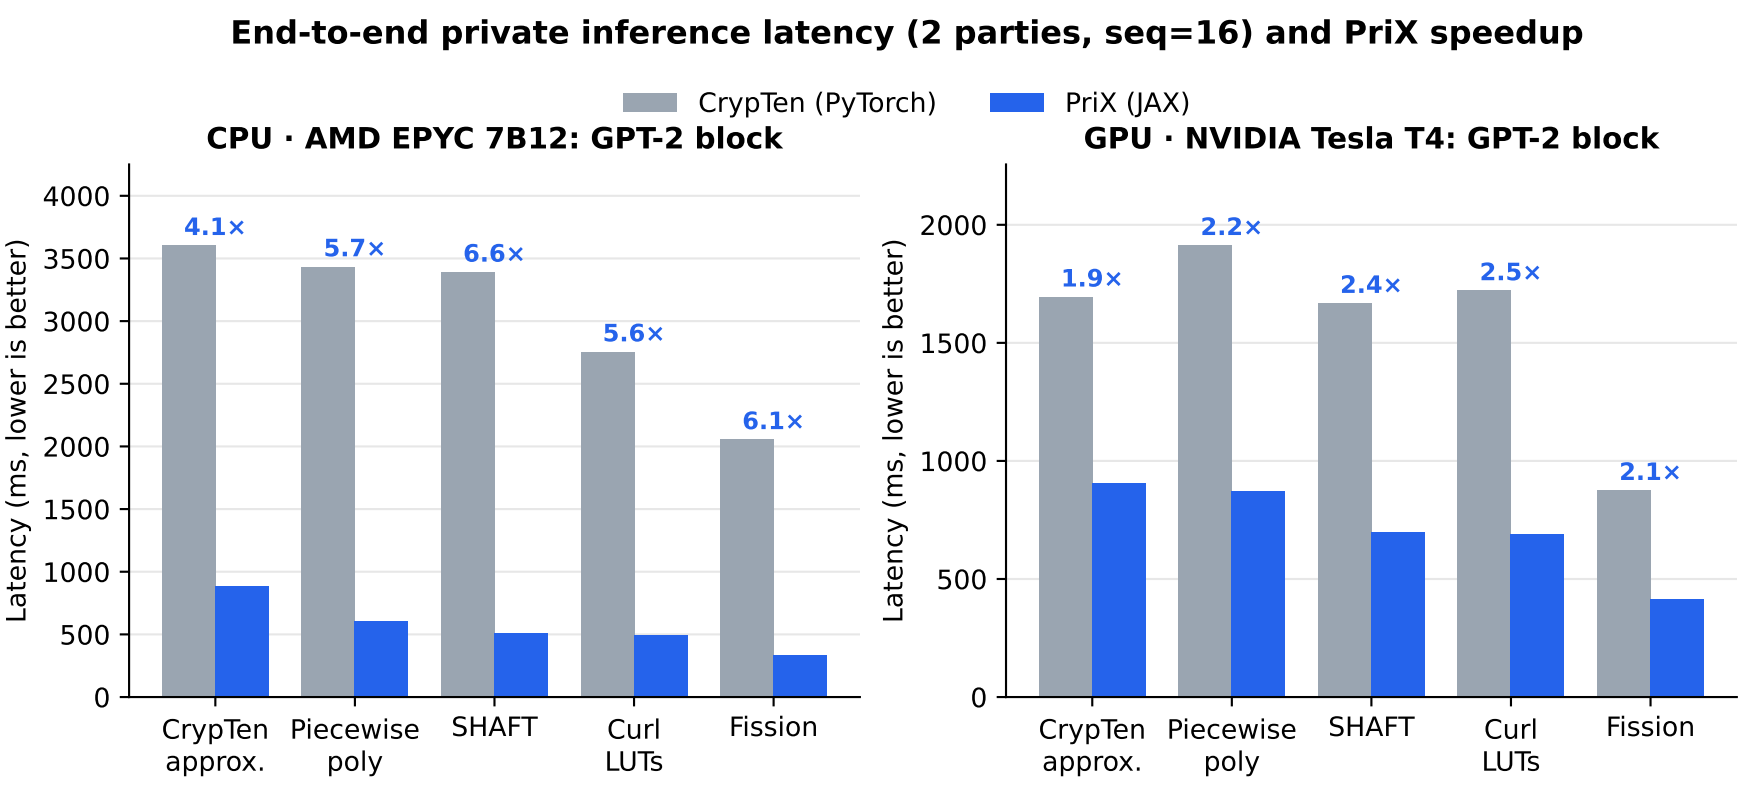
End-to-end private inference latency (2 parties, seq=16) and PriX speedup
CrypTen (PyTorch)
PriX (JAX)
CPU · AMD EPYC 7B12: GPT-2 block
GPU · NVIDIA Tesla T4: GPT-2 block
Latency (ms, lower is better)
4000
3500
3000
2500
2000
1500
1000
500
0
4.1×
5.7×
6.6×
5.6×
6.1×
1.9×
2.2×
2.4×
2.5×
2.1×
CrypTen approx.
Piecewise poly
SHAFT
Curl LUTs
Fission