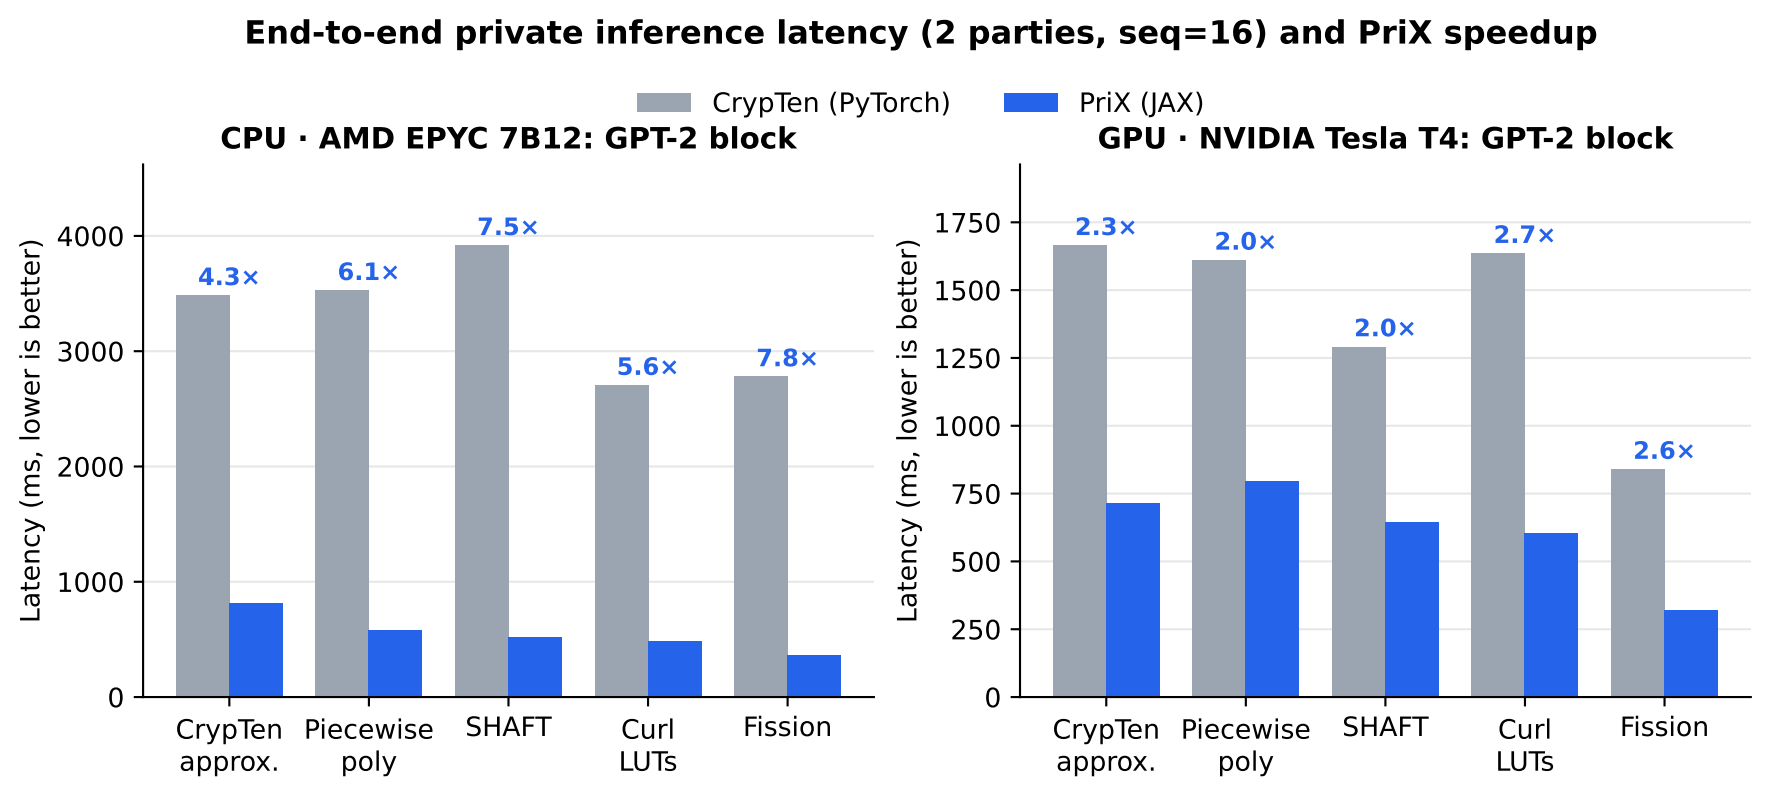
End-to-end private inference latency (2 parties, seq=16) and PriX speedup
CrypTen (PyTorch)
PriX (JAX)
CPU · AMD EPYC 7B12: GPT-2 block
GPU · NVIDIA Tesla T4: GPT-2 block
Latency (ms, lower is better)
4000
3000
2000
1000
0
4.3×
6.1×
7.5×
5.6×
7.8×
1750
1500
1250
1000
750
500
250
2.3×
2.0×
2.0×
2.7×
2.6×
CrypTen approx.
Piecewise poly
SHAFT
Curl LUTs
Fission