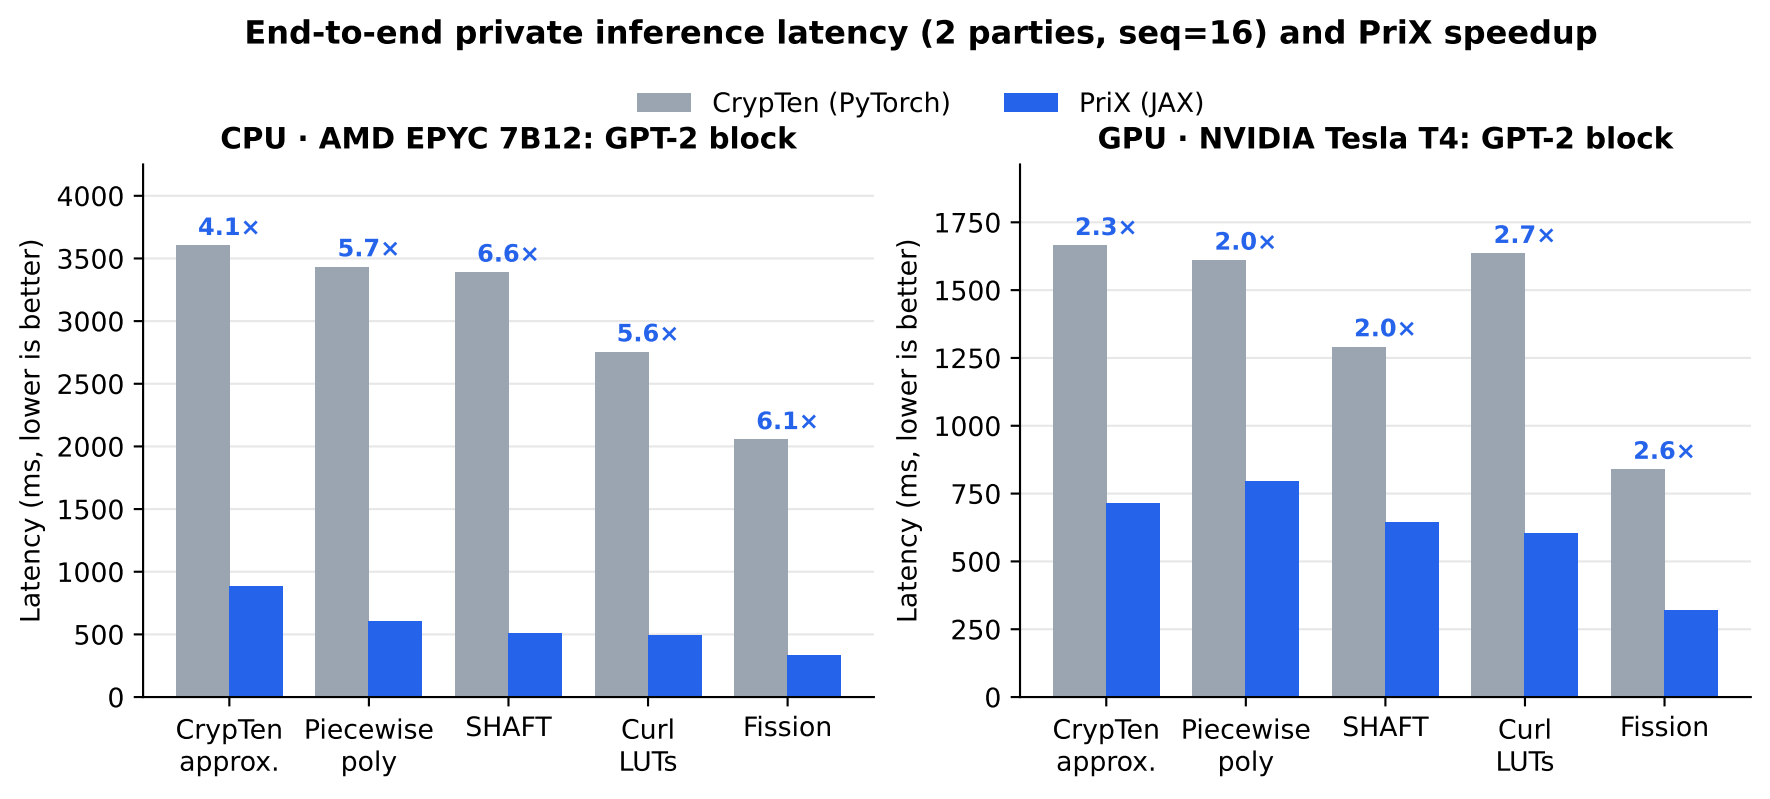
End-to-end private inference latency (2 parties, seq=16) and PriX speedup
CrypTen (PyTorch)
PriX (JAX)
CPU · AMD EPYC 7B12: GPT-2 block
GPU · NVIDIA Tesla T4: GPT-2 block
Latency (ms, lower is better)
4.1×
5.7×
6.6×
5.6×
6.1×
2.3×
2.0×
2.0×
2.7×
2.6×
CrypTen approx.
Piecewise poly
SHAFT
Curl LUTs
Fission