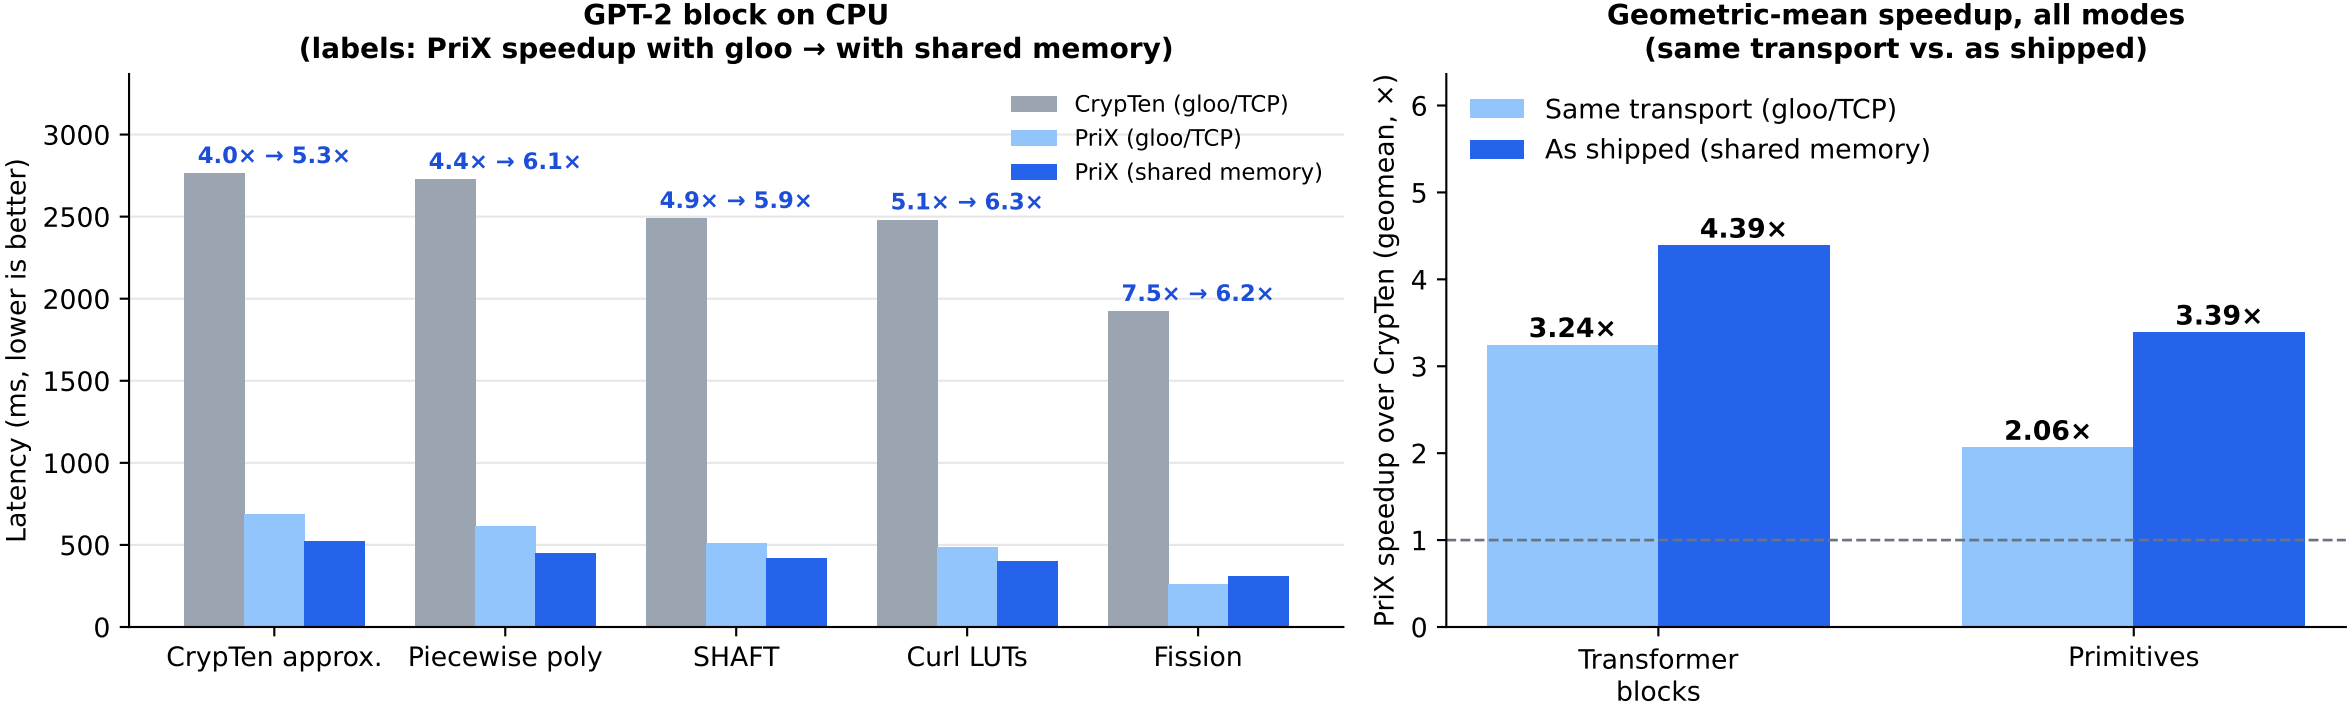
GPT-2 block on CPU
(labels: PriX speedup with gloo → with shared memory)
Latency (ms, lower is better)
CrypTen (gloo/TCP)
PriX (gloo/TCP)
PriX (shared memory)
4.0× → 5.3×
4.4× → 6.1×
4.9× → 5.9×
5.1× → 6.3×
7.5× → 6.2×
CrypTen approx.
Piecewise poly
SHAFT
Curl LUTs
Fission
Geometric-mean speedup, all modes
(same transport vs. as shipped)
PriX speedup over CrypTen (geomean, ×)
Same transport (gloo/TCP)
As shipped (shared memory)
3.24×
4.39×
2.06×
3.39×
Transformer blocks
Primitives